



BKP: An R Package for Beta Kernel Process Modeling

Jiangyan Zhao

School of Statistics, ECNU

jyzhao@sfs.ecnu.edu.cn

joint work with Kunhai Qing and Jin Xu

2026.07.24




Contents

- 1 Motivation
- 2 Methodology
- 3 Package Usage



Estimating an Unknown Probability Surface



- **Observation model**

$$y_i \mid \pi(\mathbf{x}_i) \sim \text{Binomial}(m_i, \pi(\mathbf{x}_i)).$$

- **Objective**

Estimate the underlying probability surface $\pi(\mathbf{x})$ and quantify its pointwise uncertainty.



Existing Approaches and the Remaining Gap

Logistic regression

$$\pi(\mathbf{x}) = \text{logit}^{-1}(\mathbf{x}^\top \boldsymbol{\beta})$$

Fast and interpretable
Limited nonlinear flexibility

Gaussian process classification

$$\pi(\mathbf{x}) = \text{logit}^{-1}\{f(\mathbf{x})\}, \quad f(\mathbf{x}) \sim \text{GP}$$

Flexible probability surfaces
Approximate posterior inference

Motivation

Can we retain nonlinear flexibility while performing direct, closed-form updating on the probability scale?




Contents

- 1 Motivation
- 2 Methodology
- 3 Package Usage



Beta Kernel Process: Direct Probability-Scale Updating



Key idea

Local kernel weighting + **Beta-binomial conjugacy**
 $\Rightarrow$ **flexible modeling with closed-form posterior inference.**



Closed-Form Posterior Summaries

Probability estimation

$$\hat{\pi}(\mathbf{x}) = \frac{\alpha_n(\mathbf{x})}{\alpha_n(\mathbf{x}) + \beta_n(\mathbf{x})}$$

Posterior mean of the unknown probability surface

Pointwise uncertainty

$$\text{Var}\{\pi(\mathbf{x}) \mid \mathcal{D}\} = \frac{\hat{\pi}(\mathbf{x})\{1 - \hat{\pi}(\mathbf{x})\}}{\alpha_n(\mathbf{x}) + \beta_n(\mathbf{x}) + 1}$$

$$\text{CI}_{1-\delta}(\mathbf{x}) = [Q_{\delta/2}(\mathbf{x}), Q_{1-\delta/2}(\mathbf{x})]$$

Credible intervals follow from Beta posterior quantiles

Count prediction

$$Y_{\text{new}} \mid \mathcal{D}, \mathbf{x}, m_{\text{new}} \sim \text{Beta-Binomial}\{m_{\text{new}}, \alpha_n(\mathbf{x}), \beta_n(\mathbf{x})\}.$$

Closed-form prediction for future binomial counts

Computational implication

Conditional on the kernel hyperparameters, posterior summaries and predictive distributions are available without numerical posterior approximation.



Prior Specification

Non-informative prior

$$\pi(\mathbf{x}) \sim \text{Beta}(1, 1).$$

Default uniform prior when no external information is available

Informative prior with fixed mean

$$\alpha_0(\mathbf{x}) = r_0 p_0, \quad \beta_0(\mathbf{x}) = r_0(1 - p_0)$$

$$\mathbb{E}\{\pi(\mathbf{x})\} = p_0, \quad \alpha_0(\mathbf{x}) + \beta_0(\mathbf{x}) = r_0$$

A fixed prior mean p_0 with precision controlled by r_0

Data-adaptive prior

$$p(\mathbf{x}) = \sum_{i=1}^n \frac{k(\mathbf{x}, \mathbf{x}_i)}{\sum_{i=1}^n k(\mathbf{x}, \mathbf{x}_i)} \frac{y_i}{m_i}, \quad r(\mathbf{x}) = r_0 \sum_{i=1}^n k(\mathbf{x}, \mathbf{x}_i)$$

$$\alpha_0(\mathbf{x}) = r(\mathbf{x})p(\mathbf{x}), \quad \beta_0(\mathbf{x}) = r(\mathbf{x})\{1 - p(\mathbf{x})\}.$$

The prior mean follows the local empirical proportion, while the prior precision adapts to the local kernel mass



Scaled Distance and Length-Scale Parameters

Scaled distance

$$h(\mathbf{x}, \mathbf{x}'; \boldsymbol{\theta}) = \sqrt{\sum_{j=1}^d \left(\frac{x_j - x'_j}{\theta_j} \right)^2}.$$

- Smaller θ_j : more localized borrowing and greater sensitivity along dimension j
- Larger θ_j : broader borrowing over longer distances

Length-scale structure

Isotropic: $\theta_1 = \dots = \theta_d = \theta$

Anisotropic: $\boldsymbol{\theta} = (\theta_1, \dots, \theta_d)$

An isotropic kernel uses a common length scale across dimensions, whereas an anisotropic kernel allows dimension-specific borrowing ranges

Modeling role

The length-scale parameters determine how rapidly information borrowing decreases with distance.



Kernel Functions

Implemented kernels

Gaussian: $k(h) = \exp(-h^2),$

Matérn 5/2 : $k(h) = \left(1 + \sqrt{5}h + \frac{5}{3}h^2\right) \exp(-\sqrt{5}h),$

Matérn 3/2 : $k(h) = \left(1 + \sqrt{3}h\right) \exp(-\sqrt{3}h),$

Wendland: $k(h) = (\zeta h + 1) \max\{0, 1 - h\}^\zeta, \quad \zeta = \lfloor d/2 \rfloor + 3.$

Gaussian and Matérn

Noncompact support

Every observation at a finite distance receives a positive weight, although the weight decreases with distance.

Wendland

Compact support

Observations with $h \geq 1$ receive exactly zero weight.



LOOCV-Based Length-Scale Selection

Leave-one-out quantities

- $\hat{\pi}^{-i}(\mathbf{x}_i; \boldsymbol{\theta})$: predicted probability at $\mathbf{x}_i$ after excluding observation i .
- $\tilde{\pi}_i = y_i/m_i$: observed empirical proportion.

Selection criterion

$$\hat{\boldsymbol{\theta}} = \underset{\boldsymbol{\theta}}{\operatorname{argmin}} L(\boldsymbol{\theta}; \mathcal{D}_n), \quad L \in \{\text{BS}, \text{LL}\}.$$

Brier score

$$\text{BS}(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n \{ \hat{\pi}^{-i}(\mathbf{x}_i; \boldsymbol{\theta}) - \tilde{\pi}_i \}^2.$$

Log-loss

$$\text{LL}(\boldsymbol{\theta}) = -\frac{1}{n} \sum_{i=1}^n \left[\tilde{\pi}_i \log \hat{\pi}^{-i}(\mathbf{x}_i; \boldsymbol{\theta}) + (1 - \tilde{\pi}_i) \log \{1 - \hat{\pi}^{-i}(\mathbf{x}_i; \boldsymbol{\theta})\} \right].$$

All leave-one-out predictions are obtained from one kernel matrix by setting its diagonal entries to zero.



Multi-Start Hyperparameter Optimization

Log-scale reparameterization

$$\gamma_j = \log_{10}(\theta_j), \quad \theta_j = 10^{\gamma_j}.$$

$$d_\gamma = \begin{cases} 1, & \text{isotropic,} \\ d, & \text{anisotropic.} \end{cases}$$

The transformation enforces positive length-scale parameters

Optimization workflow

- 1 Generate $n_0 = 10d_\gamma$ initial values using a Latin hypercube design
- 2 Run the derivative-free SBPLX algorithm from each initial value
- 3 Retain the solution with the smallest LOOCV loss

$$s^* = \operatorname{argmin}_{1 \leq s \leq n_0} L(\hat{\gamma}^{(s)}; \mathcal{D}_n), \quad \hat{\gamma} = \hat{\gamma}^{(s^*)}.$$

Why multi-start?

The LOOCV objective may be non-convex, so multiple initial values reduce sensitivity to a single local optimum.



Contents

- 1 Motivation
- 2 Methodology
- 3 Package Usage



fit_BKP() — Data Inputs

```
fit_BKP(X, y, m, Xbounds = NULL,  
prior = "noninformative", r0 = 2, p0 = mean(y / m),  
kernel = "gaussian", isotropic = TRUE, theta = NULL,  
loss = "brier", ess = "none", n_multi_start = NULL, n_threads = 1)
```

- The first group specifies the core **data inputs**:
 - X: an $n \times d$ numeric input matrix or data frame
 - y: a length- n vector of observed success counts
 - m: a length- n vector of corresponding binomial trial counts
 - Xbounds: an optional $d \times 2$ matrix of lower and upper bounds, used to normalize X to $[0, 1]^d$ before fitting. If NULL, X is assumed to be already normalized.



