



BKP: An R Package for Beta Kernel Process Modeling

Jiangyan Zhao

School of Statistics, ECNU

jyzhao@sfs.ecnu.edu.cn

joint work with Kunhai Qing and Jin Xu

2026.07.24




Contents

1 Motivation

2 Methodology



Estimating an Unknown Probability Surface



- **Observation model**

$$y_i \mid \pi(\mathbf{x}_i) \sim \text{Binomial}(m_i, \pi(\mathbf{x}_i)).$$

- **Objective**

Estimate the underlying probability surface $\pi(\mathbf{x})$ and quantify its pointwise uncertainty.



Existing Approaches and the Remaining Gap

Logistic regression

$$\pi(\mathbf{x}) = \text{logit}^{-1}(\mathbf{x}^\top \boldsymbol{\beta})$$

Fast and interpretable
Limited nonlinear flexibility

Gaussian process classification

$$\pi(\mathbf{x}) = \text{logit}^{-1}\{f(\mathbf{x})\}, \quad f(\mathbf{x}) \sim \text{GP}$$

Flexible probability surfaces
Approximate posterior inference

Motivation

Can we retain nonlinear flexibility while performing direct, closed-form updating on the probability scale?




Contents

- 1 Motivation
- 2 Methodology



Beta Kernel Process: Direct Probability-Scale Updating



Key idea

Local kernel weighting + **Beta-binomial conjugacy**
 ⇒ **flexible modeling with closed-form posterior inference.**



Closed-Form Posterior Summaries

Probability estimation

$$\hat{\pi}(\mathbf{x}) = \frac{\alpha_n(\mathbf{x})}{\alpha_n(\mathbf{x}) + \beta_n(\mathbf{x})}$$

Posterior mean of the unknown probability surface

Pointwise uncertainty

$$\text{Var}\{\pi(\mathbf{x}) \mid \mathcal{D}\} = \frac{\hat{\pi}(\mathbf{x})\{1 - \hat{\pi}(\mathbf{x})\}}{\alpha_n(\mathbf{x}) + \beta_n(\mathbf{x}) + 1}$$

$$\text{CI}_{1-\delta}(\mathbf{x}) = [Q_{\delta/2}(\mathbf{x}), Q_{1-\delta/2}(\mathbf{x})]$$

Credible intervals follow from Beta posterior quantiles

Count prediction

$$Y_{\text{new}} \mid \mathcal{D}, \mathbf{x}, m_{\text{new}} \sim \text{Beta-Binomial}\{m_{\text{new}}, \alpha_n(\mathbf{x}), \beta_n(\mathbf{x})\}.$$

Closed-form prediction for future binomial counts

Computational implication

Conditional on the kernel hyperparameters, posterior summaries and predictive distributions are available without numerical posterior approximation.



Prior Specification

Non-informative prior

$$\pi(\mathbf{x}) \sim \text{Beta}(1, 1).$$

Default uniform prior when no external information is available

Informative prior with fixed mean

$$\alpha_0(\mathbf{x}) = r_0 p_0, \quad \beta_0(\mathbf{x}) = r_0(1 - p_0)$$

$$\mathbb{E}\{\pi(\mathbf{x})\} = p_0, \quad \alpha_0(\mathbf{x}) + \beta_0(\mathbf{x}) = r_0$$

A fixed prior mean p_0 with precision controlled by r_0

Data-adaptive prior

$$p(\mathbf{x}) = \sum_{i=1}^n \frac{k(\mathbf{x}, \mathbf{x}_i)}{\sum_{i=1}^n k(\mathbf{x}, \mathbf{x}_i)} \frac{y_i}{m_i}, \quad r(\mathbf{x}) = r_0 \sum_{i=1}^n k(\mathbf{x}, \mathbf{x}_i)$$

$$\alpha_0(\mathbf{x}) = r(\mathbf{x})p(\mathbf{x}), \quad \beta_0(\mathbf{x}) = r(\mathbf{x})\{1 - p(\mathbf{x})\}.$$

The prior mean follows the local empirical proportion, while the prior precision adapts to the local kernel mass



Scaled Distance and Length-Scale Parameters

Scaled distance

$$h(\mathbf{x}, \mathbf{x}'; \boldsymbol{\theta}) = \sqrt{\sum_{j=1}^d \left(\frac{x_j - x'_j}{\theta_j} \right)^2}.$$

- Smaller θ_j : more local borrowing and greater sensitivity along dimension j
- Larger θ_j : smoother borrowing over longer distances

Length-scale structure

Isotropic: $\theta_1 = \dots = \theta_d = \theta$

Anisotropic: $\boldsymbol{\theta} = (\theta_1, \dots, \theta_d)$

A common length scale provides uniform smoothing, whereas dimension-specific length scales allow directional adaptation

Modeling role

The length-scale parameters determine how rapidly information borrowing decreases with distance.



Kernel Functions

Implemented kernels

Gaussian: $k(h) = \exp(-h^2),$

Matérn 5/2 : $k(h) = \left(1 + \sqrt{5}h + \frac{5}{3}h^2\right) \exp(-\sqrt{5}h),$

Matérn 3/2 : $k(h) = \left(1 + \sqrt{3}h\right) \exp(-\sqrt{3}h),$

Wendland: $k(h) = (\zeta h + 1) \max(0, 1 - h)^\zeta, \quad \zeta = \lfloor d/2 \rfloor + 3.$

Gaussian and Matérn

Noncompact support

All observations receive positive weight,
although the influence decreases with
distance.

Wendland

Compact support

Observations with $h \geq 1$ receive exactly
zero weight.



